



CS 1550

Week 11 – Lab 4

Teaching Assistant

Maher Khan

CS 1550 – Project 3 Autograder is out

- **Due:** Friday, March 29, 2019 @11:59pm
- **Late:** Sunday, March 31, 2019 @11:59pm
 - 10% reduction per late day

For Python implementation

- Add the following to top of your script:

```
#!/usr/bin/env python
```

Project 3

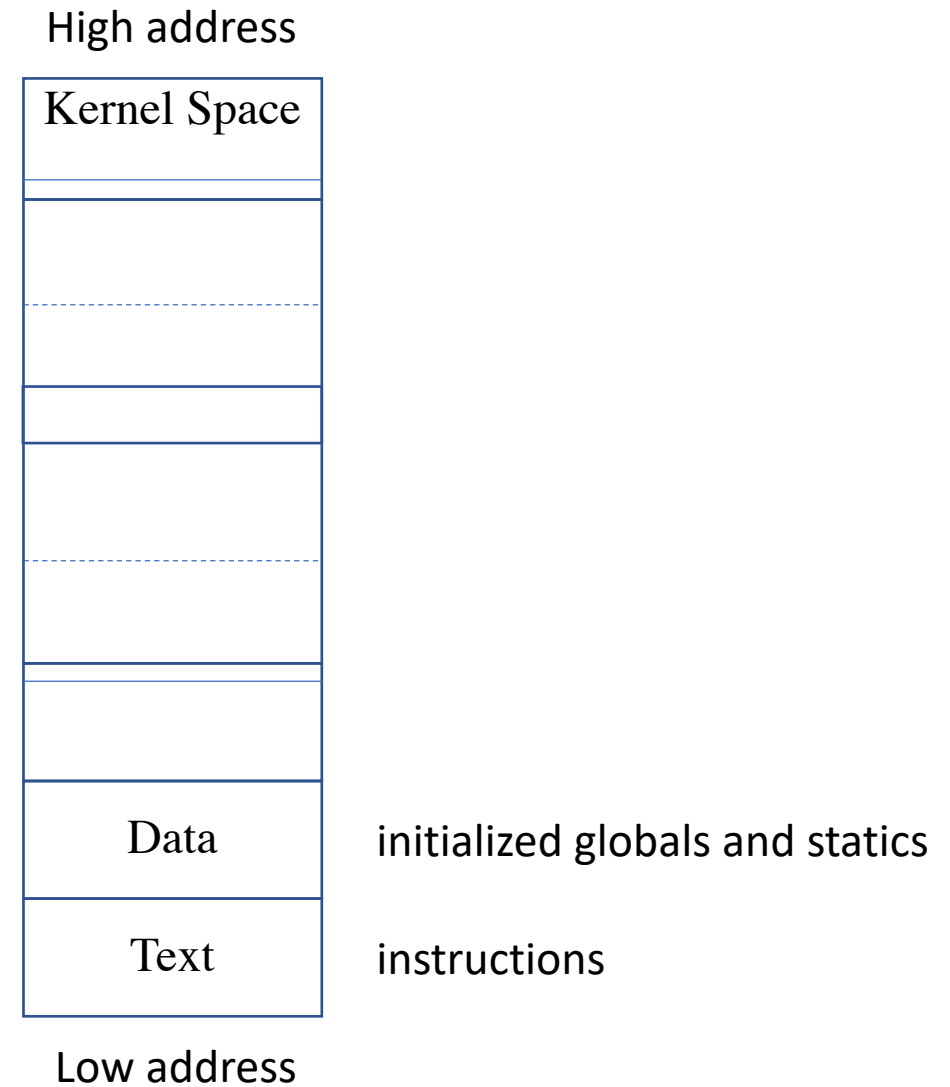
- Some implementation hints:
 - Make sure your program is efficient or it will time out
 - You only should go through the trace file once
 - Make use of data structures such as Linked List, HashMap, etc
 - Autograder may introduce very large trace files

Memory layout

```
int t = 0; // Data
...
int main() {
    ...
}
```

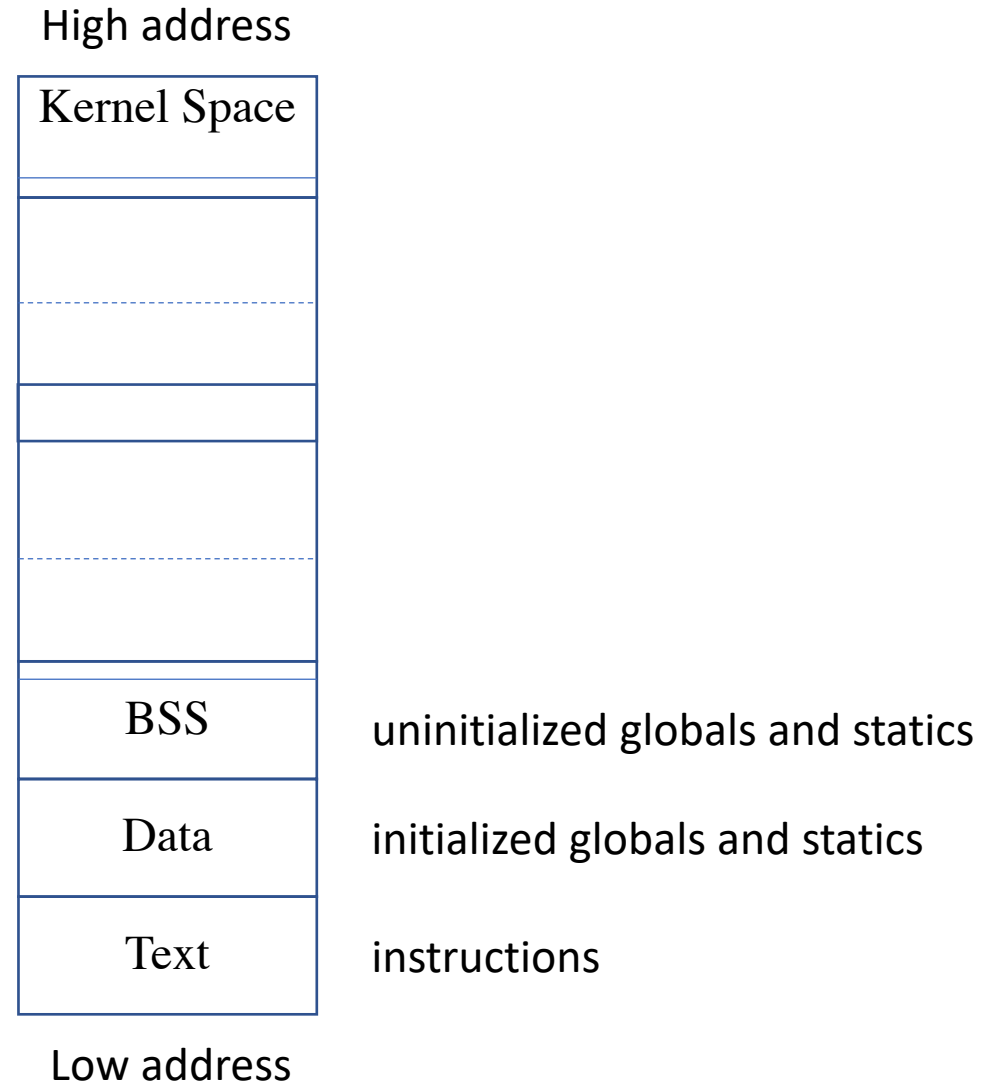
Memory layout

```
int t = 0; // Data
...
int main() {
    ...
}
```



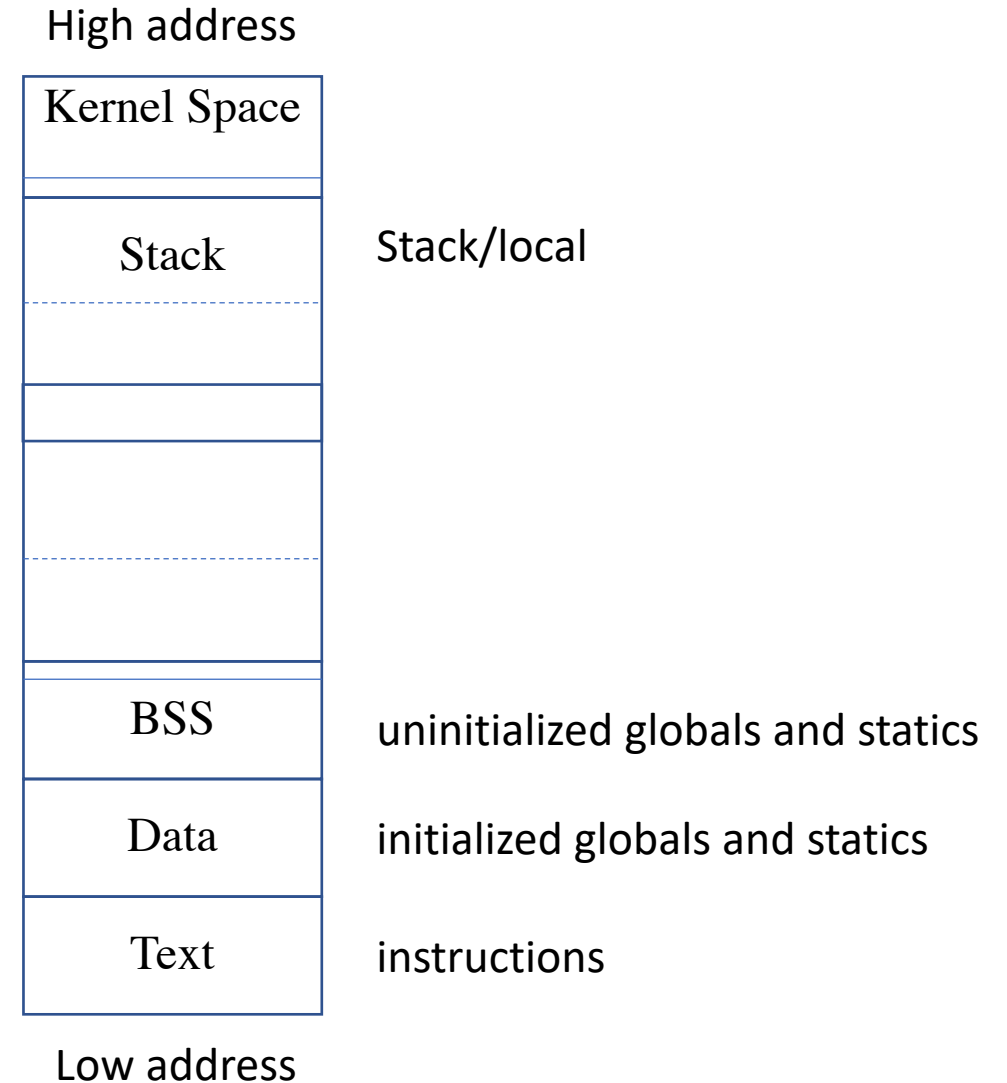
Memory layout

```
int t = 0; // Data
int m;    // BSS
...
int main() {
    ...
    static int j;    // BSS
}
```



Memory layout

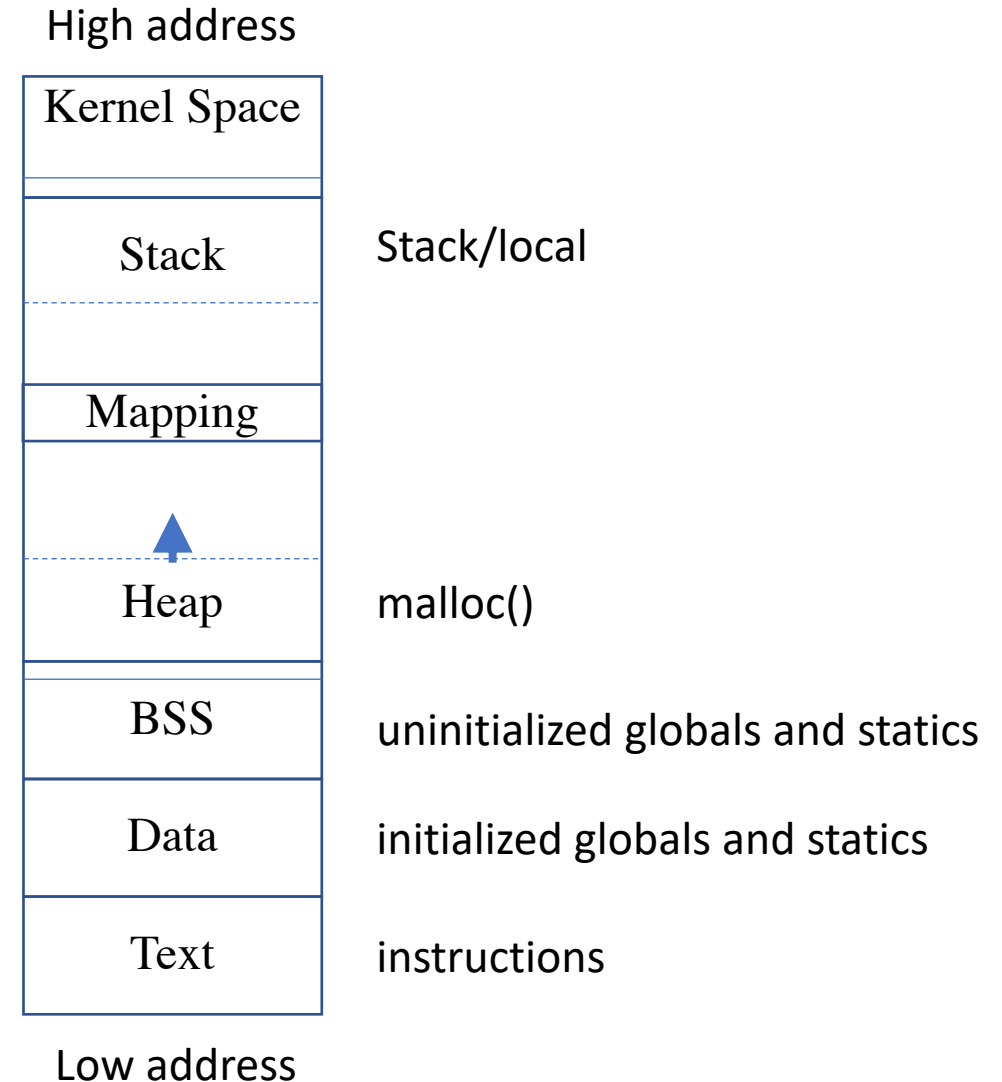
```
int t = 0; // Data
int m;    // BSS
...
int main() {
    ...
    int i;           // Stack
    static int j;    // BSS
}
```



Memory layout

```
int t = 0; // Data
int m;    // BSS
...
int main() {
    ...
    int i;           // Stack
    static int j;    // BSS

    // ptr: Stack
    // 4B pointed by ptr: Heap
    char * ptr = (char*)malloc(4);
}
```

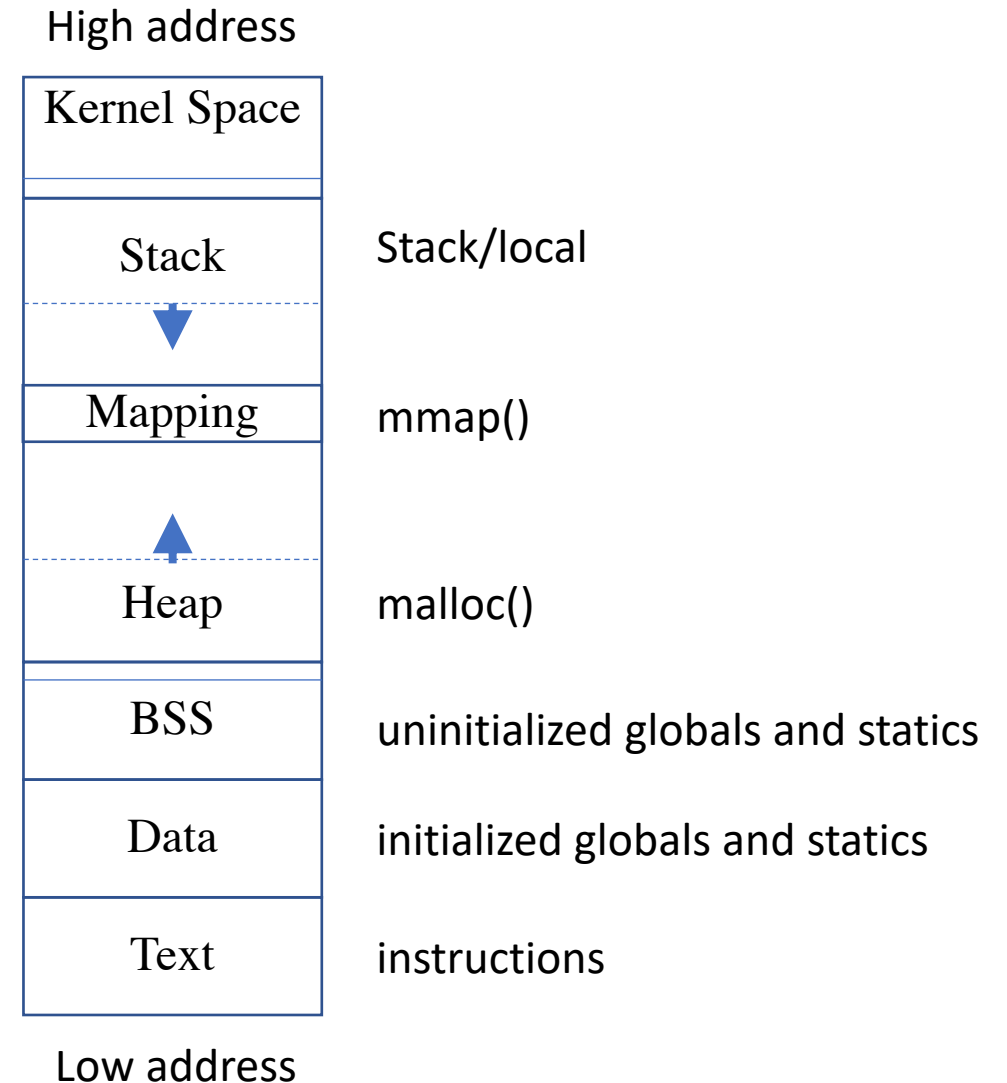


Memory layout

```
int t = 0; // Data
int m;    // BSS
...
int main() {
    ...
    int i;           // Stack
    static int j;    // BSS

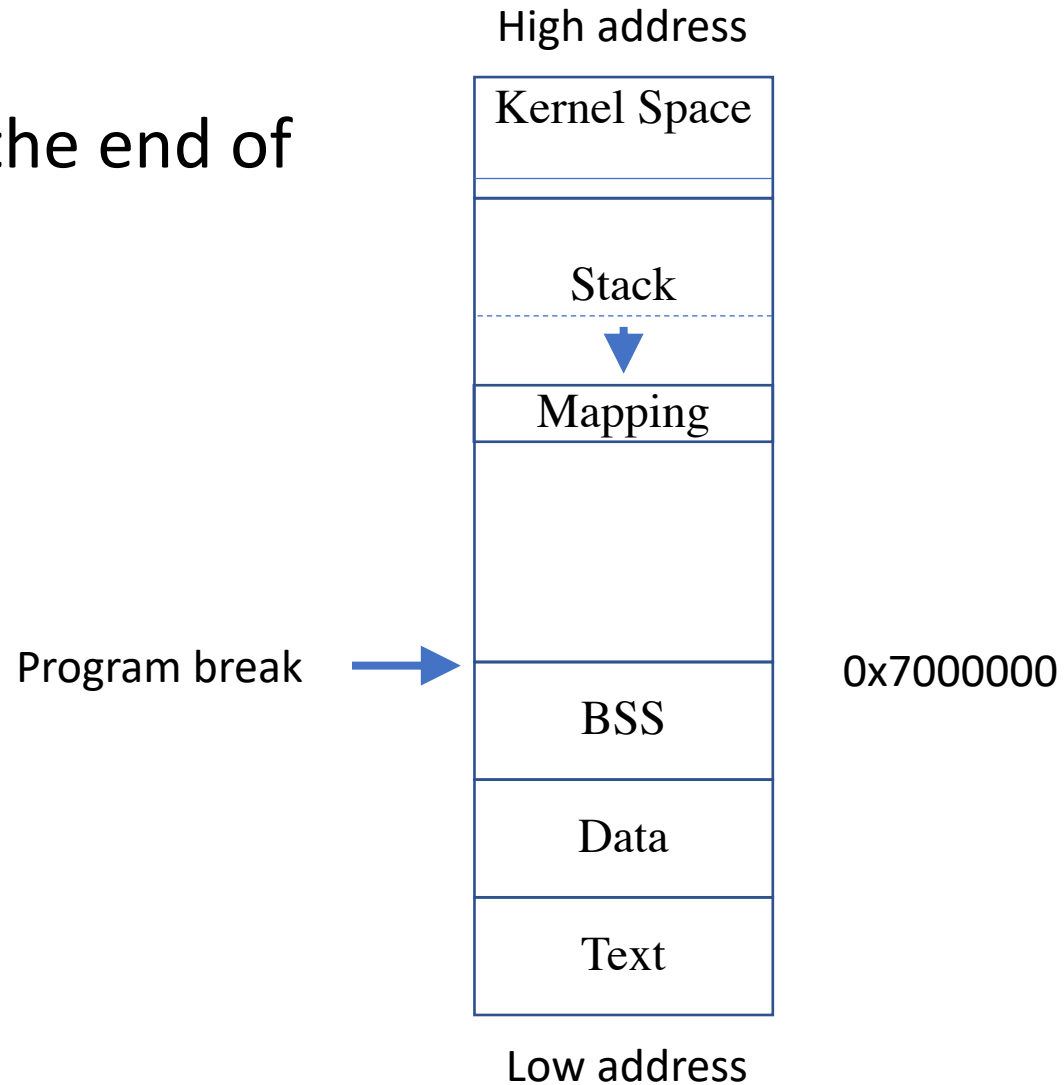
    // ptr: Stack
    // 4B pointed by ptr: Heap
    char * ptr = (char*)malloc(4);

    // mptr: Stack
    // 4K pointed by mptr: memory Mapping
    char * mptr = (char*)mmap(...,4096,...);
    ...
}
```



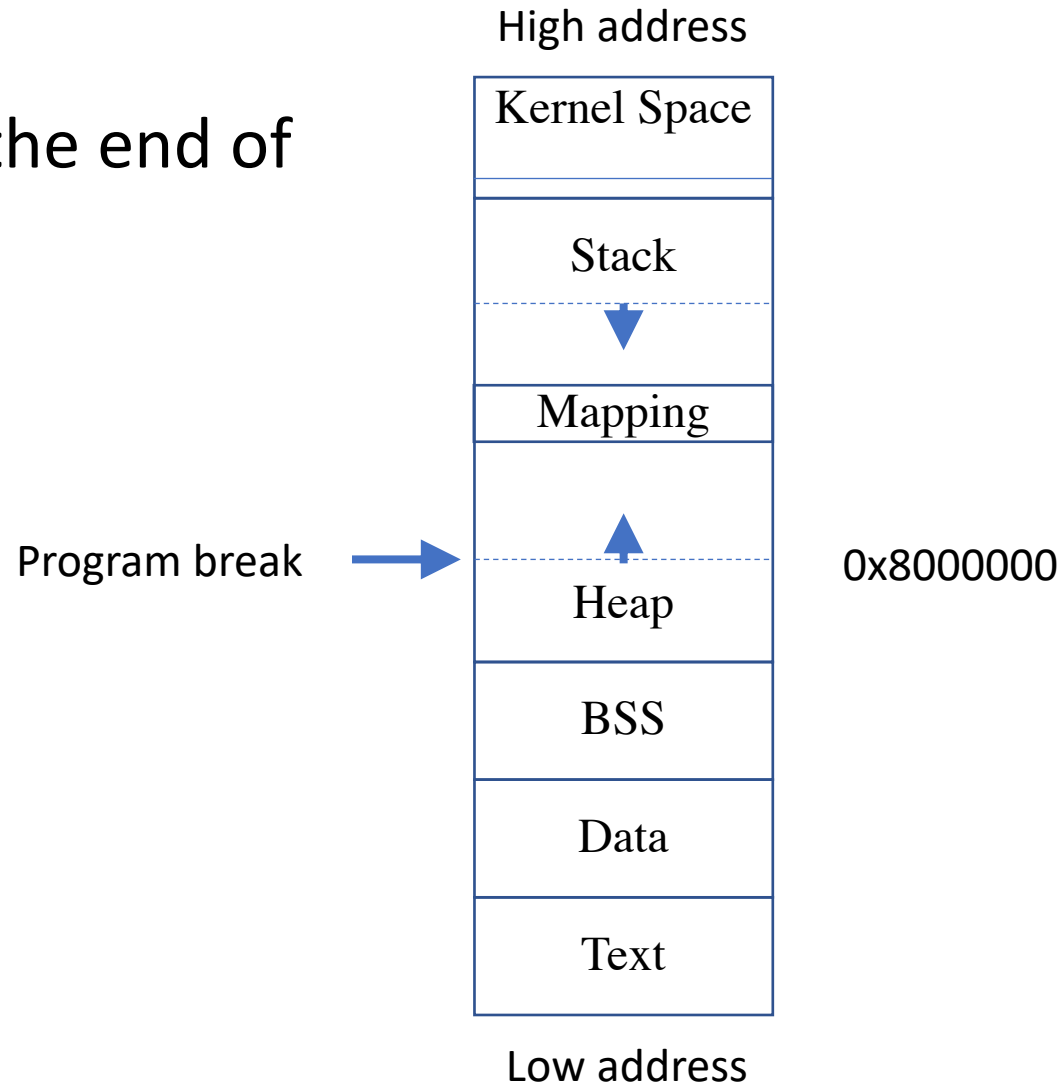
Program break

- Program break marks the end of the uninitialized data



Program break

- Program break marks the end of the uninitialized data



Program break: The syscall sbrk

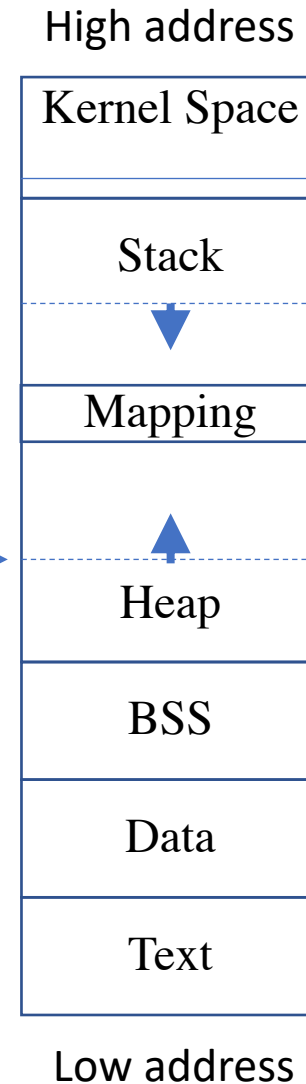
- Sbrk adds a size to the end of `cur_brk`

```
void *cur_brk = sbrk(0);
```

```
void *old_brk = sbrk(1024);
```

```
void *new_brk = sbrk(0);
```

Program break



cur_brk: 0x8000000

Program break: The syscall sbrk

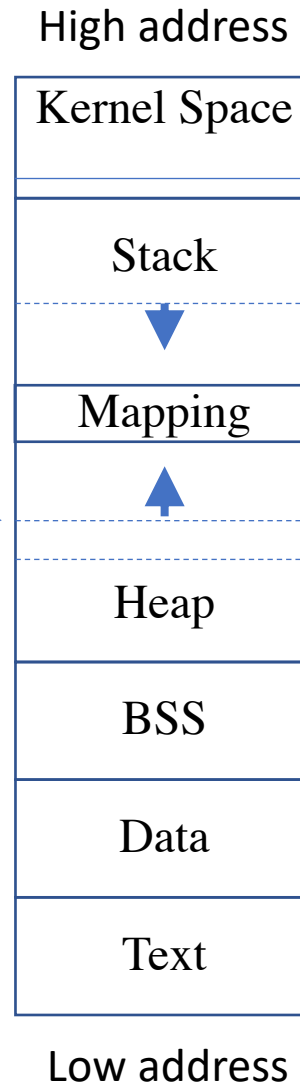
- Sbrk adds a size to the end of `cur_brk`

```
void *cur_brk = sbrk(0);
```

```
void *old_brk = sbrk(4096);
```

```
void *new_brk = sbrk(0);
```

Program break



0x8001000: increase 0x8000000 by 4K
old_brk, **cur_brk**: 0x8000000

Program break: The syscall sbrk

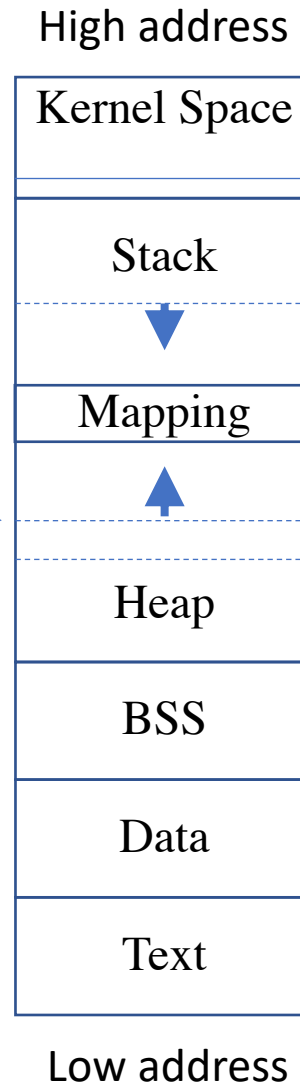
- Sbrk adds a size to the end of `cur_brk`

```
void *cur_brk = sbrk(0);
```

```
void *old_brk = sbrk(4096);
```

```
void *new_brk = sbrk(0);
```

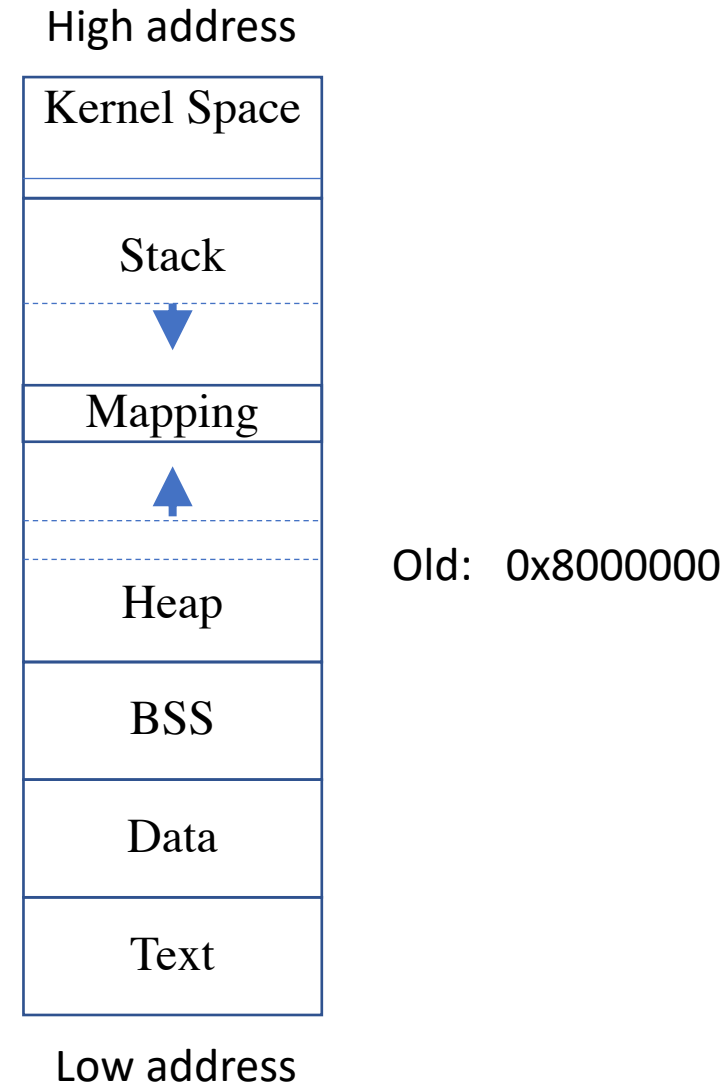
Program break



new_brk: 0x8001000
cur_brk, old_brk: 0x8000000

Program break: The syscall brk

- brk defines the absolute value for heap's end



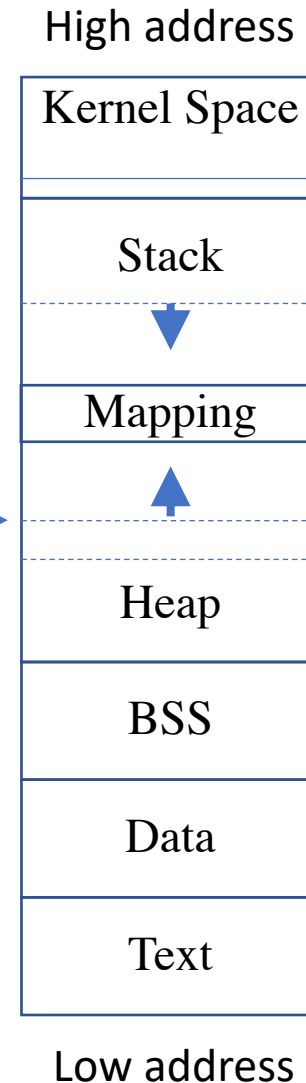
Program break: The syscall brk

- brk defines the absolute value for heap's end

```
int ret = brk(0x8001000);
```

```
if (ret != 0) {  
    // error  
}
```

Set program break to →



New: **0x8001000**
Old: 0x8000000

Sbrk on XV6

The **sys_sbrk()** in **sysproc.c** is the XV-6 implementation for sbrk.

```
...  
addr = proc->sz;  
if(growproc(n) < 0)  
    return -1;  
...  
return addr;
```

Sbrk on XV6

The **sys_sbrk()** in **sysproc.c** is the XV-6 implementation for sbrk.

```
...  
addr = proc->sz;    // get current brk  
if(growproc(n) < 0)  
    return -1;  
...  
return addr;
```

Sbrk on XV6

The **sys_sbrk()** in **sysproc.c** is the XV-6 implementation for sbrk.

...

```
addr = proc->sz;    // get current brk
if(growproc(n) < 0)    // increase brk by n
    return -1;
```

...

```
return addr;
```

growproc

The growproc() in proc.c:

...

```
if(n > 0) {  
    allocvm();  
} else if (n < 0) {  
    deallocvm();  
}
```

growproc

The growproc() in proc.c:

...

```
if(n > 0) {                                // allocation
    allocuvm();                             // allocate physical pages, update page table
} else if (n < 0) {
    deallocuvm();
}
```

growproc

The growproc() in proc.c:

...

```
if(n > 0) {           // allocation
    allocuvm();        // allocate physical pages, update page table
} else if (n < 0) {    // deallocation
    deallocuvm();      // update page table, free physical page
}
```

Physical memory allocation

Given 4KB per page and allocating an array with size of 100 pages:

```
char * ptr = (char*) malloc (4096 * 100);
```


Physical memory allocation

Given 4KB per page and allocating an array with size of 100 pages:

```
char * ptr = (char*) malloc (4096 * 100);
```

- This only allocates virtual memory: ptr to ptr+4096*100
- How about physical memory?

Physical memory allocation

Given 4KB per page and allocating an array with size of 100 pages:

```
char * ptr = (char*) malloc (4096 * 100);
```

- This only allocates virtual memory: ptr to ptr+4096*100

XV6: Immediately allocate all 100 physical page frames

Problems?

Physical memory allocation

Given 4KB per page and allocating an array with size of 100 pages:

```
char * ptr = (char*) malloc (4096 * 100);
```

```
// assume ptr is 0x8000000, i.e., a page-aligned virtual address
```

- This only allocates virtual memory: ptr to ptr+4096*100

Lab 4: allocate physical page frame upon the 1st access on that page.

```
ptr[4096*99 + 50] = 'a'; // the 1st access on the 100th page
```

Page Fault

- Page Table: Stores mapping from virtual page to physical page frame
E.g., Virtual Page 0x8000000 -> Physical 0x400000
- Translating a virtual address to physical address:
Virtual address → (TLB →) Page Table → Physical address

Page Fault

- Translating virtual address 0x8000005:
 1. Get its page-start-address 0x8000000, and **offset**-in-page **5**.
 2. Search (TLB &) Page Table to find the mapping of 0x8000000
 3. If found, e.g., 0x8000000->0x4000000:
then physical address is 0x4000000**5**.

If not
found, Page Fault

Page Fault

```
char *ptr = (char*) malloc(4096*100);
```

```
ptr[4096*99 + 50] = 'a'; // 1st access, no physical page frame: Page Fault
```

Page Fault

```
char *ptr = (char*) malloc(4096*100);  
ptr[4096*99 + 50] = 'a'; // 1st access, no physical page frame: Page Fault
```

In XV6, Page Fault on ptr[4096*99 + 50] (**inside** 100th page):

Page Fault

```
char *ptr = (char*) malloc(4096*100);  
ptr[4096*99 + 50] = 'a'; // 1st access, no physical page frame: Page Fault
```

In XV6, Page Fault on ptr[4096*99 + 50] (**inside** 100th page):

1. Issue Page Fault trap. All traps are handled by trap() in **trap.c**.

Page Fault

```
char *ptr = (char*) malloc(4096*100);  
ptr[4096*99 + 50] = 'a'; // 1st access, no physical page frame: Page Fault
```

In XV6, Page Fault on ptr[4096*99 + 50] (**inside** 100th page):

1. Issue Page Fault trap. All traps are handled by trap() in trap.c.
2. **Handle Page Fault (*Hint: T_PGFLT, how to find the faulting addr*) in trap():**
 - 1) **Allocate a physical page frame for this 100th page**
 - 2) **Update page table**

Allocate physical pages, update page table

Recall the `growproc()` in `proc.c`.

```
if(n > 0) {           // allocation
    allocvm();         // allocate physical pages, update page table
} else if (n < 0) {    // deallocation
    deallocvm();        // update page table, free physical page
}
```

Check how `allocvm` & `deallocvm` work.

(Hint: `PGROUNDDOWN`, `mappages`)

CS 1550 – Lab 4

- **Due:** Friday, 5th April, 2019 @11:59pm



CS 1550

Week 8 – Lab 4

Teaching Assistant

Maher Khan